

Import the Libraries

```
In [1]: # Import the libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import tensorflow as tf
from tensorflow.keras import layers
from tensorflow.keras.callbacks import EarlyStopping
# Remove Warnings
import warnings
warnings.filterwarnings('ignore')
```

```
2024-02-03 19:30:20.688567: E external/local_xla/xla/stream_executor/cuda/cuda_d
nn.cc:9261] Unable to register cuDNN factory: Attempting to register factory for
plugin cuDNN when one has already been registered
2024-02-03 19:30:20.688681: E external/local_xla/xla/stream_executor/cuda/cuda_
ft.cc:607] Unable to register cuFFT factory: Attempting to register factory for
plugin cuFFT when one has already been registered
2024-02-03 19:30:20.964375: E external/local_xla/xla/stream_executor/cuda/cuda_b
las.cc:1515] Unable to register cuBLAS factory: Attempting to register factory f
or plugin cuBLAS when one has already been registered
```

```
In [2]: # load the dataset
df = pd.read_csv('/kaggle/input/paddy-disease-classification/train.csv')
```

Sneak Preview of Data

Achieve a rapid glimpse of your dataset with just one line of code! This convenient approach allows you to swiftly review the initial rows of your data, providing you with an instant grasp of the dataset's content without the need to scroll through the entire file.

```
In [3]: df.head()
```

Out[3]:

	image_id	label	variety	age
0	100330.jpg	bacterial_leaf_blight	ADT45	45
1	100365.jpg	bacterial_leaf_blight	ADT45	45
2	100382.jpg	bacterial_leaf_blight	ADT45	45
3	100632.jpg	bacterial_leaf_blight	ADT45	45
4	101918.jpg	bacterial_leaf_blight	ADT45	45

```
In [4]: #Check the shape of data
print(f'The Training Dataset has {df.shape[0]} rows and {df.shape[1]} column
s.')
```

The Training Dataset has 10407 rows and 4 columns.

Let's Check Number of Unique "labels" in our Training Data

```
In [5]: df['label'].unique().tolist()
```

```
Out[5]: ['bacterial_leaf_blight',  
         'bacterial_leaf_streak',  
         'bacterial_panicle_blight',  
         'blast',  
         'brown_spot',  
         'dead_heart',  
         'downy_mildew',  
         'hispa',  
         'normal',  
         'tungro']
```

Let's Check Number of Unique "Varieties" in our Training Data

```
In [6]: df['variety'].unique().tolist()
```

```
Out[6]: ['ADT45',  
         'IR20',  
         'KarnatakaPonni',  
         'Onthanel',  
         'Ponni',  
         'Surya',  
         'Zonal',  
         'AndraPonni',  
         'AtchayaPonni',  
         'RR']
```

Observation:

1. We have 10407 images in our training data
2. We have 10 unique disease
3. We also have 10 unique varieties of rice

Summary of Data Characteristics

Descriptive statistics are employed to succinctly summarize and gain insight into the fundamental characteristics of the dataset.

In [7]: df.describe()

Out[7]:

	age
count	10407.000000
mean	64.043624
std	8.958830
min	45.000000
25%	60.000000
50%	67.000000
75%	70.000000
max	82.000000

Observation:

- We find only one numeric column **age** when we describe it:
 1. The range of age is from 45 to 82

Let's Check the values of "Age" and "Variety" by Grouping

```
In [8]: df.groupby("age")["variety"].value_counts()
```

```
Out[8]: age  variety
45  ADT45          286
    AndraPonni     112
    AtchayaPonni    55
    Surya          32
    Ponni          20
47  Ponni         112
50  ADT45         584
    KarnatakaPonni 218
    Onthanel       172
    AtchayaPonni    60
    IR20           32
55  ADT45         379
    Ponni         110
    Zonal          74
57  ADT45         126
    AtchayaPonni    87
60  ADT45        1394
    Ponni         266
62  ADT45          5
65  AndraPonni     265
    AtchayaPonni    259
    ADT45          250
66  ADT45          36
67  ADT45         401
    Ponni          14
68  ADT45         247
    Ponni           6
70  ADT45        2371
    Zonal          325
    Onthanel       179
    KarnatakaPonni  84
    IR20           82
    RR            36
72  ADT45         552
73  ADT45          38
75  KarnatakaPonni 686
    ADT45         180
77  ADT45          42
80  Ponni         129
    ADT45          96
82  ADT45           5
    Name: count, dtype: int64
```

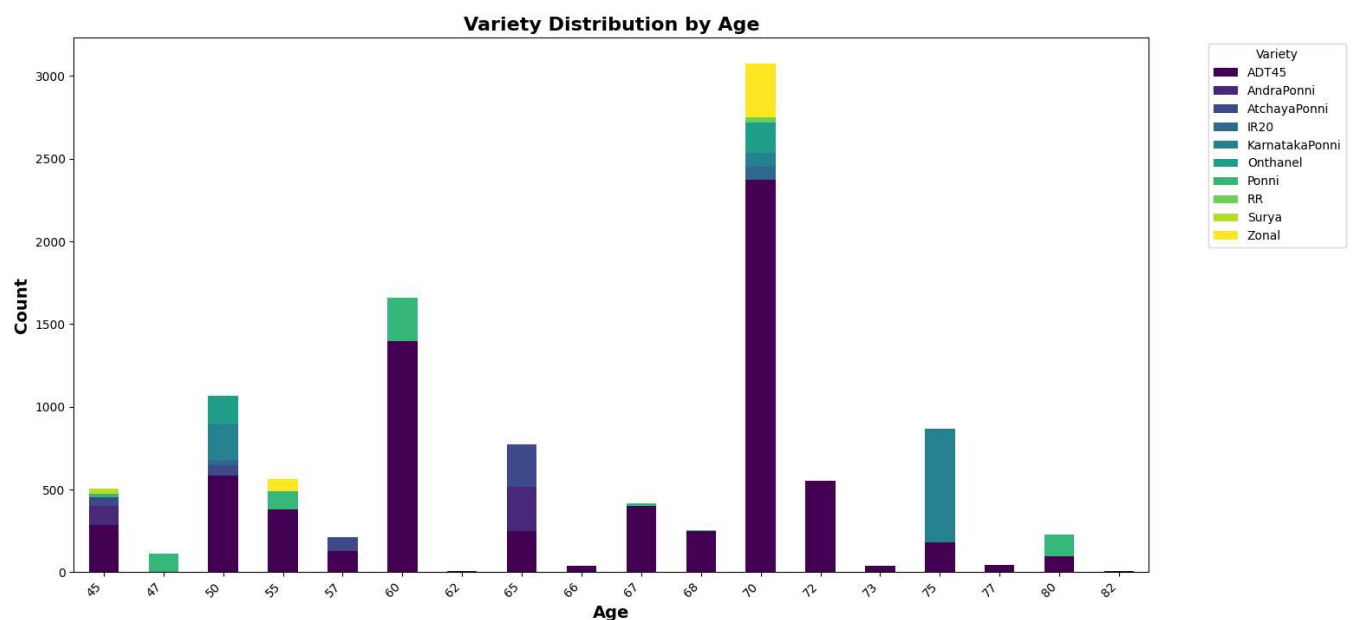
```
In [9]: # Group by age and variety, then get the counts
grouped_counts = df.groupby("age")["variety"].value_counts().unstack().fillna(0)

# Plot the grouped counts using a bar plot
fig, axes = plt.subplots(figsize=(16, 8))
grouped_counts.plot(kind='bar', stacked=True, ax=axes, cmap='viridis')

# Customize the plot
plt.title("Variety Distribution by Age", fontsize=16, fontweight='bold')
plt.xlabel("Age", fontsize=14, fontweight='bold')
plt.ylabel("Count", fontsize=14, fontweight='bold')
plt.xticks(rotation=45, ha='right')

# Add a legend
plt.legend(title='Variety', bbox_to_anchor=(1.05, 1), loc='upper left')

plt.show()
```



Let's Check the values of "Age" and "Label" by Grouping

```
In [10]: df.groupby("age")["label"].value_counts()
```

```
Out[10]: age  label
45  blast                179
    tungro               104
    hispa                 83
    downy_mildew          75
    brown_spot            48
    ...
77  dead_heart            4
80  dead_heart          129
    bacterial_panicle_blight  66
    normal                30
82  normal                 5
Name: count, Length: 95, dtype: int64
```

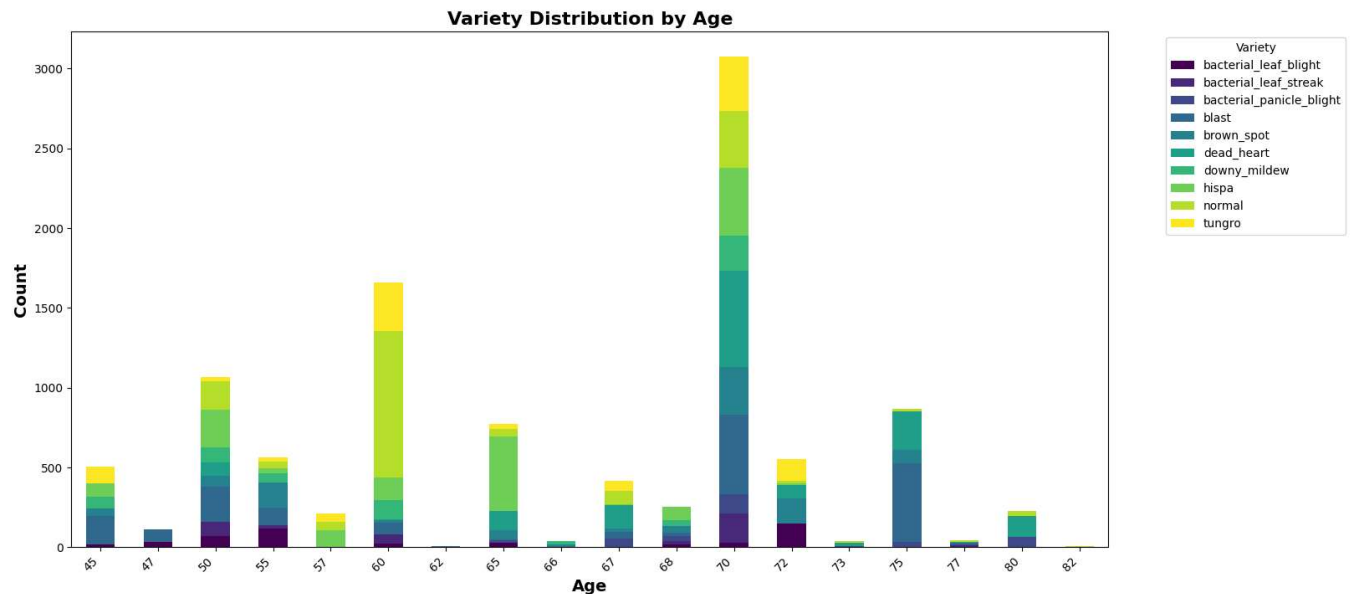
```
In [11]: # Group by age and variety, then get the counts
grouped_counts = df.groupby("age")["label"].value_counts().unstack().fillna(0)

# Plot the grouped counts using a bar plot
fig, axes = plt.subplots(figsize=(16, 8))
grouped_counts.plot(kind='bar', stacked=True, ax=axes, cmap='viridis')

# Customize the plot
plt.title("Variety Distribution by Age", fontsize=16, fontweight='bold')
plt.xlabel("Age", fontsize=14, fontweight='bold')
plt.ylabel("Count", fontsize=14, fontweight='bold')
plt.xticks(rotation=45, ha='right')

# Add a legend
plt.legend(title='Variety', bbox_to_anchor=(1.05, 1), loc='upper left')

plt.show()
```



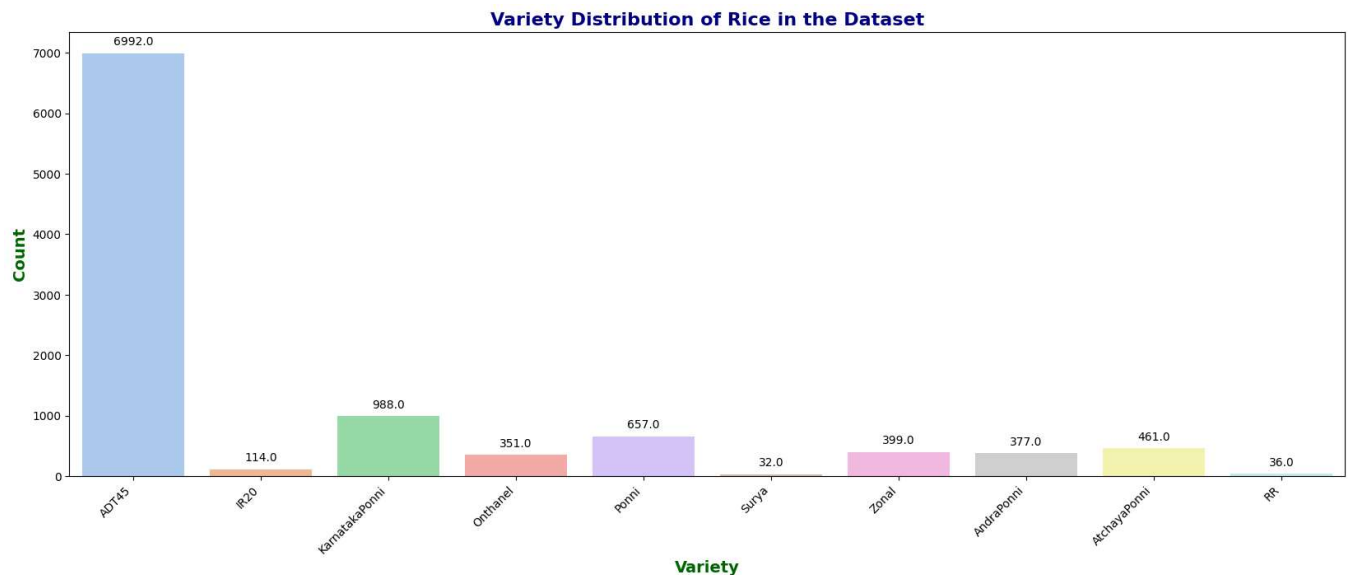
```
In [12]: # Set a custom color palette
custom_palette = sns.color_palette("pastel")

# Plot the data count based on the variety name using countplot
fig, axes = plt.subplots(1, 1, figsize=(20, 7))
sns.countplot(x='variety', data=df, palette=custom_palette, ax=axes)

# Customize the plot
plt.title("Variety Distribution of Rice in the Dataset", fontsize=16, fontweight='bold', color='navy')
plt.xlabel("Variety", fontsize=14, fontweight='bold', color='darkgreen')
plt.ylabel("Count", fontsize=14, fontweight='bold', color='darkgreen')
plt.xticks(rotation=45, ha='right') # Rotate x-axis labels for better visibility

# Add annotations (counts) on top of each bar
for p in axes.patches:
    axes.annotate(f'{p.get_height()}', (p.get_x() + p.get_width() / 2., p.get_height()),
                  ha='center', va='center', xytext=(0, 10), textcoords='offset points',
                  fontsize=10, color='black')

plt.show()
```



- ADT45 is the most common variety inside the training data

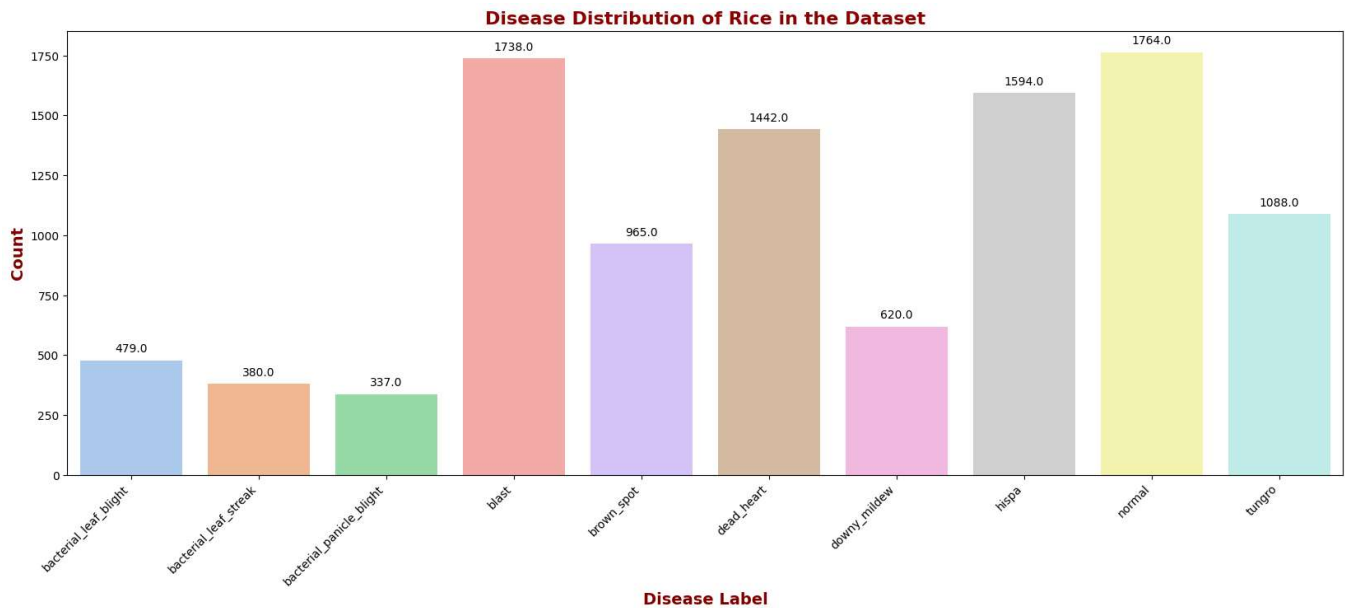
```
In [13]: # Set a custom color palette
custom_palette = sns.color_palette("pastel")

# Plot the data count based on the label using countplot
fig, axes = plt.subplots(1, 1, figsize=(20, 7))
sns.countplot(x='label', data=df, palette=custom_palette, ax=axes)

# Customize the plot
plt.title("Disease Distribution of Rice in the Dataset", fontsize=16, fontweight='bold', color='darkred')
plt.xlabel("Disease Label", fontsize=14, fontweight='bold', color='maroon')
plt.ylabel("Count", fontsize=14, fontweight='bold', color='maroon')
plt.xticks(rotation=45, ha='right') # Rotate x-axis labels for better visibility

# Add annotations (counts) on top of each bar
for p in axes.patches:
    axes.annotate(f'{p.get_height()}', (p.get_x() + p.get_width() / 2., p.get_height()),
                  ha='center', va='center', xytext=(0, 10), textcoords='offset points',
                  fontsize=10, color='black')

plt.show()
```



```
In [14]: normal = df[df['label']=='normal']
normal = normal[normal['variety']=='ADT45']
normal_items = normal.image_id[:7].values
normal_items.tolist()
```

```
Out[14]: ['100007.jpg',
'100025.jpg',
'100135.jpg',
'100165.jpg',
'100171.jpg',
'100186.jpg',
'100188.jpg']
```

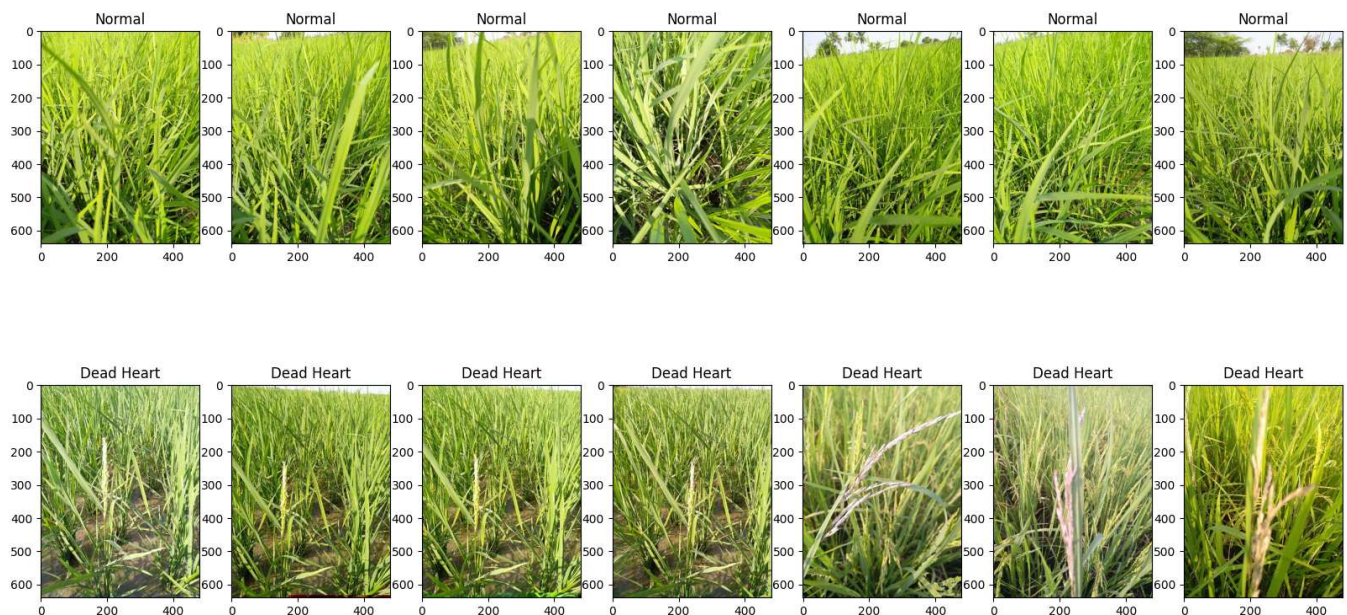


```
In [15]: dead = df[df['label']=='dead_heart']
dead = dead[dead['variety']=='ADT45']
dead_items = dead.image_id[:7].values
dead_items.tolist()
```

```
Out[15]: ['101165.jpg',
'102750.jpg',
'108367.jpg',
'109900.jpg',
'100222.jpg',
'100670.jpg',
'100715.jpg']
```

```
In [16]: plt.figure(figsize=(20, 10))
columns = 7
path = '/kaggle/input/paddy-disease-classification/train_images/'

for i, image_loc in enumerate(np.concatenate((normal_items, dead_items))):
    plt.subplot(10//columns + 1, columns, i+1)
    if i < 7:
        image = plt.imread(path+"normal/"+image_loc)
        plt.title('Normal')
    else:
        image = plt.imread(path+"dead_heart/"+image_loc)
        plt.title('Dead Heart')
    plt.imshow(image)
```



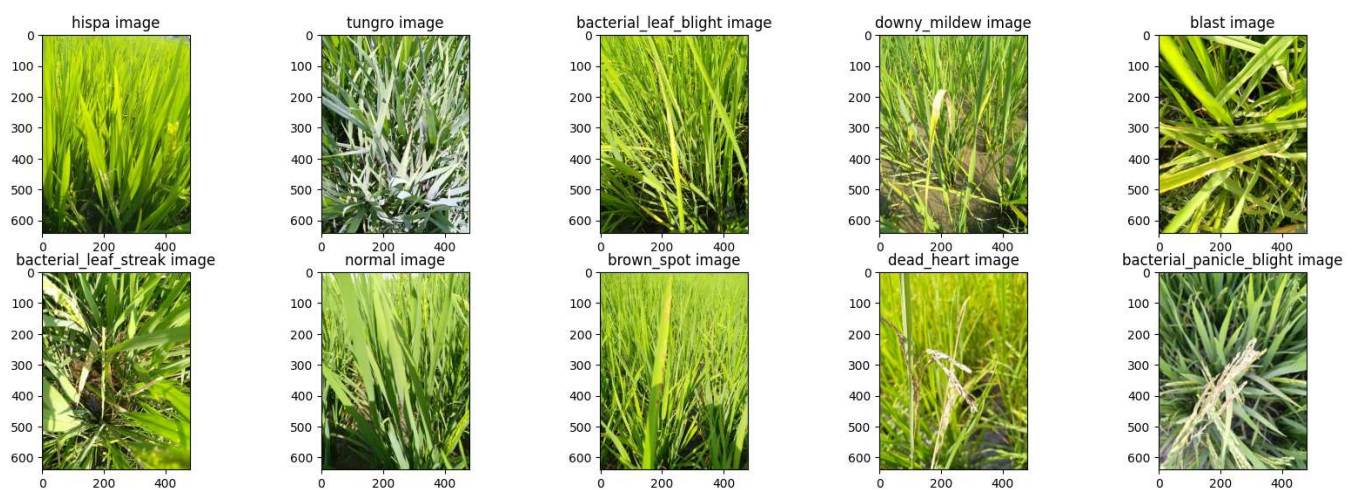
```

In [17]: images = [
            '/kaggle/input/paddy-disease-classification/train_images/hispa/106590.jpg',
            '/kaggle/input/paddy-disease-classification/train_images/tungro/109629.jpg',
            '/kaggle/input/paddy-disease-classification/train_images/bacterial_leaf_blight/109372.jpg',
            '/kaggle/input/paddy-disease-classification/train_images/downy_mildew/102350.jpg',
            '/kaggle/input/paddy-disease-classification/train_images/blast/110243.jpg',
            '/kaggle/input/paddy-disease-classification/train_images/bacterial_leaf_streak/101104.jpg',
            '/kaggle/input/paddy-disease-classification/train_images/normal/109760.jpg',
            '/kaggle/input/paddy-disease-classification/train_images/brown_spot/104675.jpg',
            '/kaggle/input/paddy-disease-classification/train_images/dead_heart/105159.jpg',
            '/kaggle/input/paddy-disease-classification/train_images/bacterial_panicle_blight/101351.jpg'
        ]

diseases = ['hispa' , 'tungro', 'bacterial_leaf_blight', 'downy_mildew', 'blast',
            "bacterial_leaf_streak" , 'normal' , 'brown_spot' , 'dead_heart' ,
            'bacterial_panicle_blight' ]

diseases = [disease + ' image' for disease in diseases]
plt.figure(figsize=(20,10))
columns = 5
for i, image_loc in enumerate(images):
    plt.subplot(len(images)//columns+1,columns,i+1)
    image = plt.imread(image_loc)
    plt.title(diseases[i])
    plt.imshow(image)

```



```

In [18]: from sklearn.preprocessing import LabelEncoder
encoder = LabelEncoder()

df['label'] = encoder.fit_transform(df['label'])
df['variety'] = encoder.fit_transform(df['variety'])

```

```
In [19]: df.head()
```

```
Out[19]:
```

	image_id	label	variety	age
0	100330.jpg	0	0	45
1	100365.jpg	0	0	45
2	100382.jpg	0	0	45
3	100632.jpg	0	0	45
4	101918.jpg	0	0	45

```
In [20]: batch_size=32  
image_width=224  
image_height=224
```

```
In [21]: train_ds = tf.keras.utils.image_dataset_from_directory(  
    directory=path,  
    validation_split = 0.3,  
    subset = "training",  
    seed = 123,  
    image_size = (image_height, image_width),  
    batch_size = batch_size  
)
```

Found 10407 files belonging to 10 classes.
Using 7285 files for training.

```
In [22]: val_ds = tf.keras.utils.image_dataset_from_directory(  
    directory=path,  
    validation_split=0.2,  
    subset="validation",  
    seed=123,  
    image_size=(image_height, image_height),  
    batch_size=batch_size  
)
```

Found 10407 files belonging to 10 classes.
Using 2081 files for validation.

```
In [23]: class_names = train_ds.class_names  
print(class_names)
```

```
['bacterial_leaf_blight', 'bacterial_leaf_streak', 'bacterial_panicle_blight',  
'blast', 'brown_spot', 'dead_heart', 'downy_mildew', 'hispa', 'normal', 'tungro']
```

```
In [24]: for image_batch, label_batch in train_ds:  
    print(image_batch.shape)  
    print(label_batch.shape)  
    break
```

```
(32, 224, 224, 3)  
(32,)
```

```
In [25]: normalization_layer = tf.keras.layers.Rescaling(1./255)
```

```
In [26]: normalized_ds = train_ds.map(lambda x,y: (normalization_layer(x),y))
image_batch , label_batch = next(iter(normalized_ds))
first_image = image_batch[0]
# Notice the pixel values are now in `[0,1]`
print(np.min(first_image),np.max(first_image))
```

```
0.0 1.0
```

```
In [27]: num_classes = len(class_names)

model = tf.keras.Sequential([
    tf.keras.layers.Rescaling(1./255),
    tf.keras.layers.Conv2D(32, 3, activation='relu'),
    tf.keras.layers.MaxPooling2D(),
    tf.keras.layers.Conv2D(64, 3, activation='relu'),
    tf.keras.layers.MaxPooling2D(),
    tf.keras.layers.Conv2D(128, 3, activation='relu'),
    tf.keras.layers.MaxPooling2D(),
    tf.keras.layers.Conv2D(256, 3, activation='relu'),
    tf.keras.layers.MaxPooling2D(),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dense(64, activation='relu'),
    tf.keras.layers.Dropout(0.15),
    tf.keras.layers.Dense(num_classes, activation='softmax')
])
```

```
In [28]: model.compile(
    optimizer='adam',
    loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
    metrics=['accuracy']
)
```

```
In [29]: %%time
# Define the callback function
early_stopping = EarlyStopping(patience=20)

history = model.fit(train_ds,
                    validation_data = val_ds,
                    epochs=100,
                    callbacks=[early_stopping])

# Evaluate the model
loss = model.evaluate(val_ds)

# Plotting the training and testing loss
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper right')
plt.show()

# Plotting the training and testing accuracy
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='lower left')
```


Epoch 1/100

WARNING: All log messages before absl::InitializeLog() is called are written to STDERR

I0000 00:00:1706988665.162202 81 device_compiler.h:186] Compiled cluster using XLA! This line is logged at most once for the lifetime of the process.


228/228 [=====] - 34s 110ms/step - loss: 1.9249 - accuracy: 0.3224 - val_loss: 1.6389 - val_accuracy: 0.4435
Epoch 2/100
228/228 [=====] - 15s 65ms/step - loss: 1.3959 - accuracy: 0.5198 - val_loss: 1.1937 - val_accuracy: 0.6290
Epoch 3/100
228/228 [=====] - 16s 67ms/step - loss: 0.8987 - accuracy: 0.7072 - val_loss: 0.9953 - val_accuracy: 0.7232
Epoch 4/100
228/228 [=====] - 15s 65ms/step - loss: 0.5132 - accuracy: 0.8350 - val_loss: 0.9019 - val_accuracy: 0.7727
Epoch 5/100
228/228 [=====] - 16s 67ms/step - loss: 0.2664 - accuracy: 0.9108 - val_loss: 0.9803 - val_accuracy: 0.7943
Epoch 6/100
228/228 [=====] - 15s 65ms/step - loss: 0.1495 - accuracy: 0.9546 - val_loss: 1.0909 - val_accuracy: 0.7866
Epoch 7/100
228/228 [=====] - 15s 66ms/step - loss: 0.1258 - accuracy: 0.9596 - val_loss: 1.1720 - val_accuracy: 0.7991
Epoch 8/100
228/228 [=====] - 15s 66ms/step - loss: 0.0876 - accuracy: 0.9757 - val_loss: 1.1111 - val_accuracy: 0.8011
Epoch 9/100
228/228 [=====] - 16s 66ms/step - loss: 0.0714 - accuracy: 0.9790 - val_loss: 1.2641 - val_accuracy: 0.7871
Epoch 10/100
228/228 [=====] - 16s 66ms/step - loss: 0.0678 - accuracy: 0.9784 - val_loss: 1.0509 - val_accuracy: 0.8059
Epoch 11/100
228/228 [=====] - 16s 67ms/step - loss: 0.0449 - accuracy: 0.9864 - val_loss: 1.2164 - val_accuracy: 0.8030
Epoch 12/100
228/228 [=====] - 15s 66ms/step - loss: 0.0662 - accuracy: 0.9809 - val_loss: 1.1235 - val_accuracy: 0.8155
Epoch 13/100
228/228 [=====] - 16s 67ms/step - loss: 0.0445 - accuracy: 0.9868 - val_loss: 1.1316 - val_accuracy: 0.8251
Epoch 14/100
228/228 [=====] - 15s 66ms/step - loss: 0.0283 - accuracy: 0.9916 - val_loss: 1.3199 - val_accuracy: 0.8270
Epoch 15/100
228/228 [=====] - 16s 68ms/step - loss: 0.0525 - accuracy: 0.9850 - val_loss: 1.2800 - val_accuracy: 0.8011
Epoch 16/100
228/228 [=====] - 15s 66ms/step - loss: 0.0326 - accuracy: 0.9900 - val_loss: 1.2694 - val_accuracy: 0.8198
Epoch 17/100
228/228 [=====] - 16s 68ms/step - loss: 0.0350 - accuracy: 0.9903 - val_loss: 1.1777 - val_accuracy: 0.8087
Epoch 18/100
228/228 [=====] - 15s 66ms/step - loss: 0.0465 - accuracy: 0.9857 - val_loss: 1.2410 - val_accuracy: 0.8188
Epoch 19/100
228/228 [=====] - 16s 67ms/step - loss: 0.0218 - accuracy: 0.9934 - val_loss: 1.4057 - val_accuracy: 0.8169
Epoch 20/100
228/228 [=====] - 15s 66ms/step - loss: 0.0384 - accuracy: 0.9886 - val_loss: 1.4273 - val_accuracy: 0.8217
Epoch 21/100
228/228 [=====] - 16s 68ms/step - loss: 0.0287 - accuracy:


```
cy: 0.9919 - val_loss: 1.3381 - val_accuracy: 0.8227
Epoch 22/100
228/228 [=====] - 15s 66ms/step - loss: 0.0237 - accuracy: 0.9929 - val_loss: 1.5807 - val_accuracy: 0.7991
Epoch 23/100
228/228 [=====] - 16s 67ms/step - loss: 0.0452 - accuracy: 0.9856 - val_loss: 1.1707 - val_accuracy: 0.8126
Epoch 24/100
228/228 [=====] - 15s 65ms/step - loss: 0.0377 - accuracy: 0.9896 - val_loss: 1.2614 - val_accuracy: 0.8136
66/66 [=====] - 3s 40ms/step - loss: 1.2614 - accuracy: 0.8136
```



CPU times: user 17min, sys: 50.8 s, total: 17min 51s
Wall time: 6min 39s

Out[29]: <matplotlib.legend.Legend at 0x7cb6d44a69b0>



```
In [30]: model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
rescaling_1 (Rescaling)	(None, 224, 224, 3)	0
conv2d (Conv2D)	(None, 222, 222, 32)	896
max_pooling2d (MaxPooling2D)	(None, 111, 111, 32)	0
conv2d_1 (Conv2D)	(None, 109, 109, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 54, 54, 64)	0
conv2d_2 (Conv2D)	(None, 52, 52, 128)	73856
max_pooling2d_2 (MaxPooling2D)	(None, 26, 26, 128)	0
conv2d_3 (Conv2D)	(None, 24, 24, 256)	295168
max_pooling2d_3 (MaxPooling2D)	(None, 12, 12, 256)	0
flatten (Flatten)	(None, 36864)	0
dense (Dense)	(None, 512)	18874880
dense_1 (Dense)	(None, 64)	32832
dropout (Dropout)	(None, 64)	0
dense_2 (Dense)	(None, 10)	650

=====
Total params: 19296778 (73.61 MB)
Trainable params: 19296778 (73.61 MB)
Non-trainable params: 0 (0.00 Byte)
=====

```
In [31]: loss , accu = model.evaluate(val_ds)
print(f"the Testing loss is {loss:.2f}")
print(f"The testing accuracy is {accu*100:.2f}%")
```

```
66/66 [=====] - 3s 41ms/step - loss: 1.2614 - accuracy: 0.8136
the Testing loss is 1.26
The testing accuracy is 81.36%
```

```
In [32]: test_ds = tf.keras.utils.image_dataset_from_directory(
        '/kaggle/input/paddy-disease-classification/test_images/',
        image_size=(image_height, image_width),
        batch_size=batch_size,
        label_mode=None,
        shuffle=False)
```

Found 3469 files belonging to 1 classes.

```
In [33]: y_pred = model.predict(test_ds, batch_size = batch_size, verbose = 1)
        y_pred.shape
```

109/109 [=====] - 9s 79ms/step

Out[33]: (3469, 10)

```
In [34]: y_pred_classes = y_pred.argmax(axis = 1)
        y_pred_classes.shape
```

Out[34]: (3469,)

```
In [35]: y_classes_names = [class_names[x] for x in y_pred_classes]
```

```
In [36]: predictions = pd.read_csv('/kaggle/input/paddy-disease-classification/sample_submission.csv')
        predictions['label'] = y_classes_names
        predictions.to_csv('submission.csv', index = False)
```

```
In [ ]:
```